

Roberto Echeverria Sosa

(323) 809-5834 | recheverriasosa@ucsd.edu | www.linkedin.com/in/robertoecheverriasosa | <https://github.com/Robert-ech>

EDUCATION

University of California San Diego

Bachelor of Science in Data Science, Machine Learning Track

La Jolla, CA

Expected Dec 2027

Relevant Coursework: Data Management (SQL & Databases), Machine Learning & Representation Learning, Data Structures & Algorithms

TECHNICAL SKILLS

Languages: Python, R, SQL, JavaScript/TypeScript, HTML/CSS

Libraries: Pandas, NumPy, Scikit-learn, PyTorch, XGBoost, PySpark, statsmodels, React, D3.js, FastAPI

Platforms & Tools: Databricks (Azure), Delta Lake, ADLS Gen2, Streamlit, Tableau, Power BI, Git, Agentic Development

WORK EXPERIENCE

Sumitomo Mitsui Banking Corporation (SMBC)

New Jersey/New York

Enterprise Architecture Intern

Jun 2026 - Aug 2026

- Co-designed and deployed AI document intelligence system (Databricks, Claude RAG, React/FastAPI): 90% reduction in analyst review time (~\$0.05/doc); partnered with credit teams on requirements, built domain-specific extraction logic for structured schema output, validated across 140+ production documents, and iterated on bugs and edge cases to achieve 95%+ accuracy
- Developed extraction platform standard with Databricks teams: built 6-step RAG pipeline (Claude via Databricks Model Serving, Delta Lake) with adaptive extraction methods for small/medium/large documents, achieving 95%+ field accuracy; adopted as the foundation for all document extraction across SMBC's AI app portfolio
- Pitched 5-year AI transformation to executives: 12-engineer team building 800+ production apps to drive adoption 200% and boost ROTE by 15%; later consulted on competitive positioning vs. Tier 1 banks during Q&A

Data Science Student Society at UCSD

San Diego, CA

Consulting Director

Jan 2026 - Present

- Established strategic partnerships between the club and companies including JPMorgan for consulting projects, pitching slide decks and presenting directly to management teams
- Built interactive dashboard for CER client analyzing 16K+ IEA records; developed ARIMA forecasts showing \$53B economic upside from 1.5X EV adoption across top 10 oil importers by 2030

PROJECTS

Azure Databricks ETL Pipeline (PySpark, Delta Lake)

- Built an end-to-end ETL pipeline on Azure Databricks processing 1M+ bank fraud records with PySpark and Delta Lake on ADLS Gen2, implementing medallion architecture (Bronze → Silver → Gold) with schema enforcement and data quality flagging
- Reduced raw data volume ~40% through deduplication and quality filtering in PySpark transformation notebooks

Credit Risk Predictor (Python, XGBoost, Streamlit)

- Cleaned 2M Lending Club loans (151 features), eliminating 60+ data leakage columns; built XGBoost credit default model achieving ROC AUC = 0.735 and 68% default recall on 20% base rate
- Deployed interactive Streamlit dashboard with real-time risk scoring, SHAP explainability, and dynamic waterfall visualizations

Fantasy Football Draft App (Gradient Boosting, NFL Data API, React)

- Processed 16,000+ rows of NFL data covering 3 seasons and 300+ players via automated API pipelines
- Built custom Fantasy Point Projection system using Gradient Boosting Regressor (MAE = 0.52) with full-stack draft interface featuring interactive player stats and highlight video playback

LEADERSHIP AND PROFESSIONAL DEVELOPMENT

Management Leadership for Tomorrow

Jan 2026 - Present

Career Preparation Fellow

- Selected for 18-month professional development program accelerating career growth through structured coaching, mentorship, and exposure to partner organizations including LinkedIn, Bloomberg, and Deloitte